

STA402L: HOMEWORK 2

DUE: 11:59 PM ON FRIDAY, JANUARY 30

Instructions. Solutions must be submitted to Gradescope as a single PDF. Programming exercises must be completed in R, should be clearly presented, and include all R code. Lab questions are restated here for convenience, but you should refer to the lab itself for details.

Total points. Book exercises: 16; Lab exercises 14; Overall: 30.

BOOK EXERCISES

B1. (3 points) Show that the posterior variance of the beta–binomial model can be written as

$$\text{Var}(\theta \mid y) = \frac{\mathbb{E}(\theta \mid y) \mathbb{E}(1 - \theta \mid y)}{a + b + n + 1}.$$

B2. Hoff 3.1.

- (a) (1 point)
- (b) (1 point)
- (c) (1 point)
- (d) (1 point)
- (e) (2 points)

B3. Hoff 3.3.

- (a) (3 points)
- (b) (3 points)
- (c) (1 point)

LAB EXERCISES

L1. (2 points) Plot a histogram of θ from the `rstan` object called `pool_output`. Describe the distribution.

- L2. (2 points) Visualize the posterior distributions of the θ_i with boxplots. In the plot, there should be one box and whiskers object for each θ_i .
- L3. (2 points) Take a few minutes to look at the contents of the two files `lab-02-pool.stan` and `lab-02-nopool.stan`. How are they different?
- L4. (2 points) What observable quantity does the parameter a represent about our prior beliefs with respect to these data? What does b represent?
- L5. (2 points) What do we actually observe in the rat tumor data with respect to these quantities?
- L6. (2 points) How well do our different prior beliefs—the ones represented by the different parameter settings above—match up with the data?
- L7. (2 points) Why might we have observed such a difference between the two approaches when using the prior $\text{Beta}(1, 1)$?